



# CS 4495 Computer Vision

## *Linear Filtering 1: Filters, Convolution, Smoothing*

---



Aaron Bobick  
School of Interactive Computing



# Linear outline

- Images are really functions  $\vec{I}(x, y)$  where the vector can be any dimension but typical are 1, 3, and 4. (When 4?) Or thought of as a multi-dimensional *signal* as a function of spatial location.
- **Image processing** is (mostly) computing new functions of image functions. Many involve **linear operators**.
- Very useful linear operator is *convolution / correlation* - what most people call filtering – because the new value is determined by local values.
- With convolution can do things like noise reduction, smoothing, and edge finding (last one is next time).

# Images as functions



# Images as functions

- We can think of an image as a function,  $f$  or  $I$ , from  $\mathbb{R}^2$  to  $\mathbb{R}$ :

$f(x, y)$  gives the *intensity* or value at position  $(x, y)$

Realistically, we expect the image only to be defined over a rectangle, with a finite range:



$$f: [a, b] \times [c, d] \rightarrow [0, 1.0] \text{ (why sometimes 255???)}$$

- A color image is just three functions “pasted” together. We can write this as a “vector-valued” function:

$$f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$$

# Digital images



- In computer vision we typically operate on **digital (discrete)** images:
  - **Sample** the 2D space on a regular grid
  - **Quantize** each sample (round to “nearest integer”)
- Image thus represented as a matrix of integer values.



$j$  →

62	79	23	119	120	105	4	0
10	10	9	62	12	78	34	0
10	58	197	46	46	0	0	48
176	135	5	188	191	68	0	49
2	1	1	29	26	37	0	77
0	89	144	147	187	102	62	208
255	252	0	166	123	62	0	31
166	63	127	17	1	0	99	30



2D



1D

# Matlab – images are matrices

```
>> im = imread('peppers.png'); % semicolon or many numbers  
>> imgreen = im(:, :, 2);  
>> imshow(imgreen)  
>> line([1 512], [256 256], 'color', 'r')  
>> plot(imgreen(256, :));
```



# Noise in images

- Noise as an example of images really being functions
- Noise is just another function that is combined with the original function to get a new – guess what – function

$$\vec{I}'(x, y) = \vec{I}(x, y) + \vec{\eta}(x, y)$$

- In images noise looks, well, noisy.

# Common types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- ***Gaussian noise:*** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise




Impulse noise



Gaussian noise



# Gaussian noise



$$f(x, y) = \underbrace{\hat{f}(x, y)}_{\text{Ideal Image}} + \underbrace{\eta(x, y)}_{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:  
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

```
>> noise = randn(size(im)).*sigma;
```

```
>> output = im + noise;
```


# Effect of $\sigma$ on Gaussian noise

**Image shows  
the noise  
values  
themselves.**


```
noise = randn(size(im)).*sigma;
```




Sigma = 2



Sigma = 8



Sigma = 32



Sigma = 64

# BE VERY CAREFUL!!!

- In previous slides, I did not say (at least wasn't supposed to say) what the range of the image was. A  $\sigma$  of 1.0 would be tiny if the range is [0 255] but huge if [0.0 1.0].
- Matlab can do either and you need to be very careful. If in doubt convert to double.
- Even more difficult can be displaying the image. Things like:
  - `imshow(I, [LOW HIGH])`  
display the image from [low high]

Don't worry – you'll get used to these hassles... see problem set PS0.

# Back to our program...


# Suppose want to remove the noise...

# First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Assumptions:
  - Expect pixels to be like their neighbors
  - Expect noise processes to be independent from pixel to pixel


# First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Moving average in 1D:



# Weighted Moving Average


- Can add weights to our moving average
- *Weights*  $[1, 1, 1, 1, 1] / 5$





# Weighted Moving Average

- Non-uniform weights [1, 4, 6, 4, 1] / 16



# Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0								

# Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10							

# Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20						

# Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30					

# Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30				

# Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

# Correlation filtering

Say the averaging window size is  $2k+1 \times 2k+1$ :

$$G[i, j] = \underbrace{\frac{1}{(2k+1)^2}}_{\text{Attribute uniform weight to each pixel}} \underbrace{\sum_{u=-k}^k \sum_{v=-k}^k F[i+u, j+v]}_{\text{Loop over all pixels in neighborhood around image pixel } F[i,j]}$$

*Attribute uniform weight to each pixel*      *Loop over all pixels in neighborhood around image pixel  $F[i,j]$*

Now generalize to allow **different weights** depending on neighboring pixel's relative position:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k \underbrace{H[u, v]}_{\text{Non-uniform weights}} F[i+u, j+v]$$



# Correlation filtering

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$


This is called **cross-correlation**, denoted  $G = H \otimes F$

Filtering an image: replace each pixel with a linear combination of its neighbors.

The filter “**kernel**” or “**mask**”  $H[u, v]$  is the prescription for the weights in the linear combination.

# Averaging filter

- What values belong in the kernel  $H$  for the moving average example?



$$G = H \otimes F$$

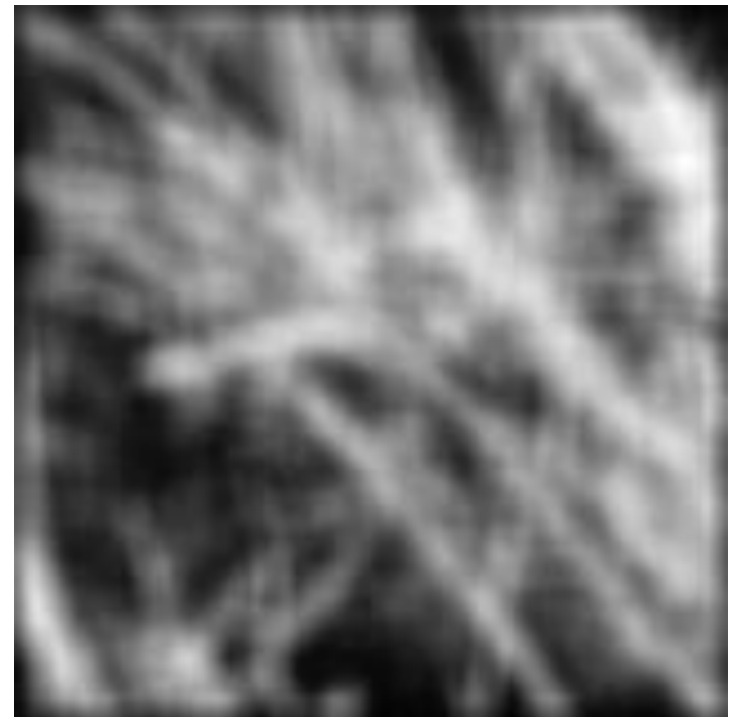
# Smoothing by averaging



depicts box filter:  
white = high value, black = low value



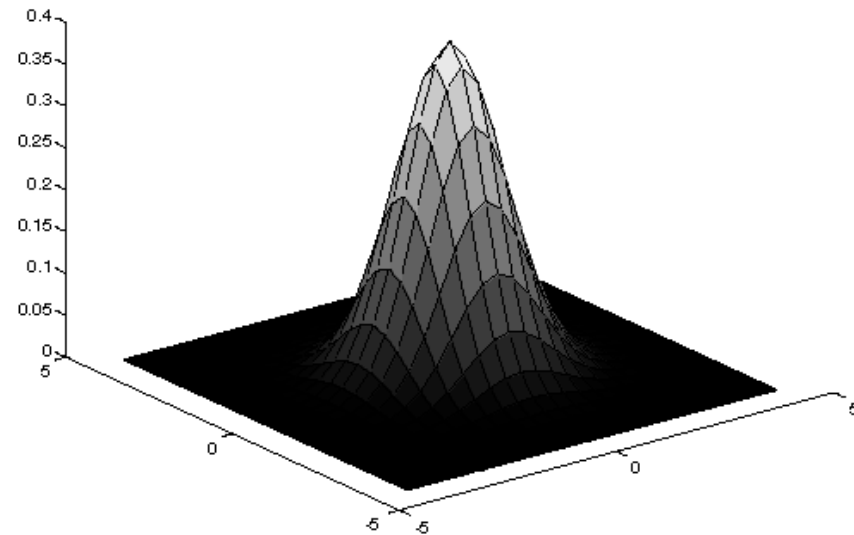
**original**



**filtered**

# Squares aren't smooth...

- Smoothing with an average actually doesn't compare at all well with a defocussed lens
- Most obvious difference is that a single point of light viewed in a defocussed lens looks like a fuzzy blob; but the averaging process would give a little square.
- More about “impulse” responses later...



# Gaussian filter

- What if we want nearest neighboring pixels to have the most influence on the output?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$F[x, y]$


$\frac{1}{16}$

1	2	1
2	4	2
1	2	1


$H[u, v]$

This kernel is an approximation of a Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$




# An Isotropic Gaussian





The picture shows a smoothing kernel proportional to

$$\exp\left(-\frac{(x^2 + y^2)}{2\sigma^2}\right)$$


(which is a reasonable model of a circularly symmetric fuzzy blob)



# Smoothing with a Gaussian



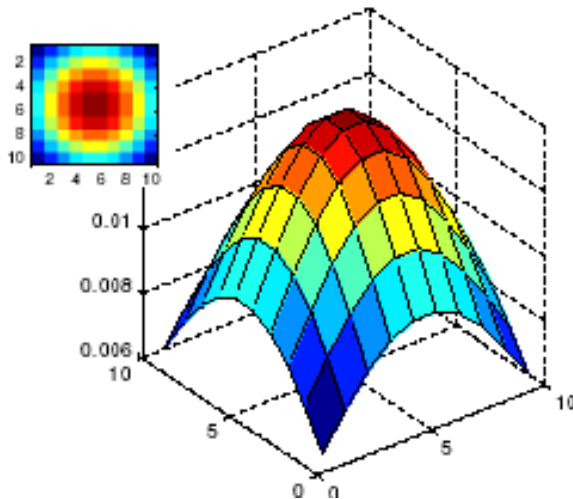
# Smoothing with not a Gaussian



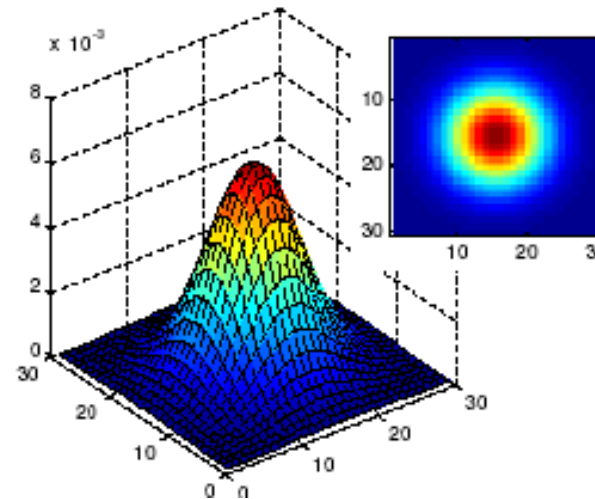


# Gaussian filters

- What parameters matter here?
- **Size** of kernel or mask
  - Note, Gaussian function has infinite support, but discrete filters use finite kernels




$\sigma = 5$  with  
 $10 \times 10$   
kernel



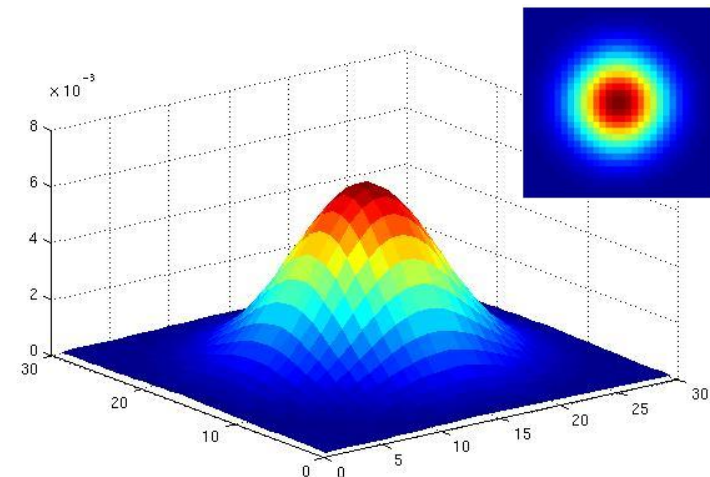
$\sigma = 5$  with  
 $30 \times 30$   
kernel

# Gaussian filters

- What parameters matter here?
- **Variance** of Gaussian: determines extent of smoothing



$\sigma = 2$  with  
 $30 \times 30$   
kernel




$\sigma = 5$  with  
 $30 \times 30$   
kernel

# Matlab

```
>> hsize = 10;  
>> sigma = 5;  
>> h = fspecial('gaussian', hsize, sigma);
```

```
>> mesh(h) ;
```



```
>> imagesc(h) ;
```



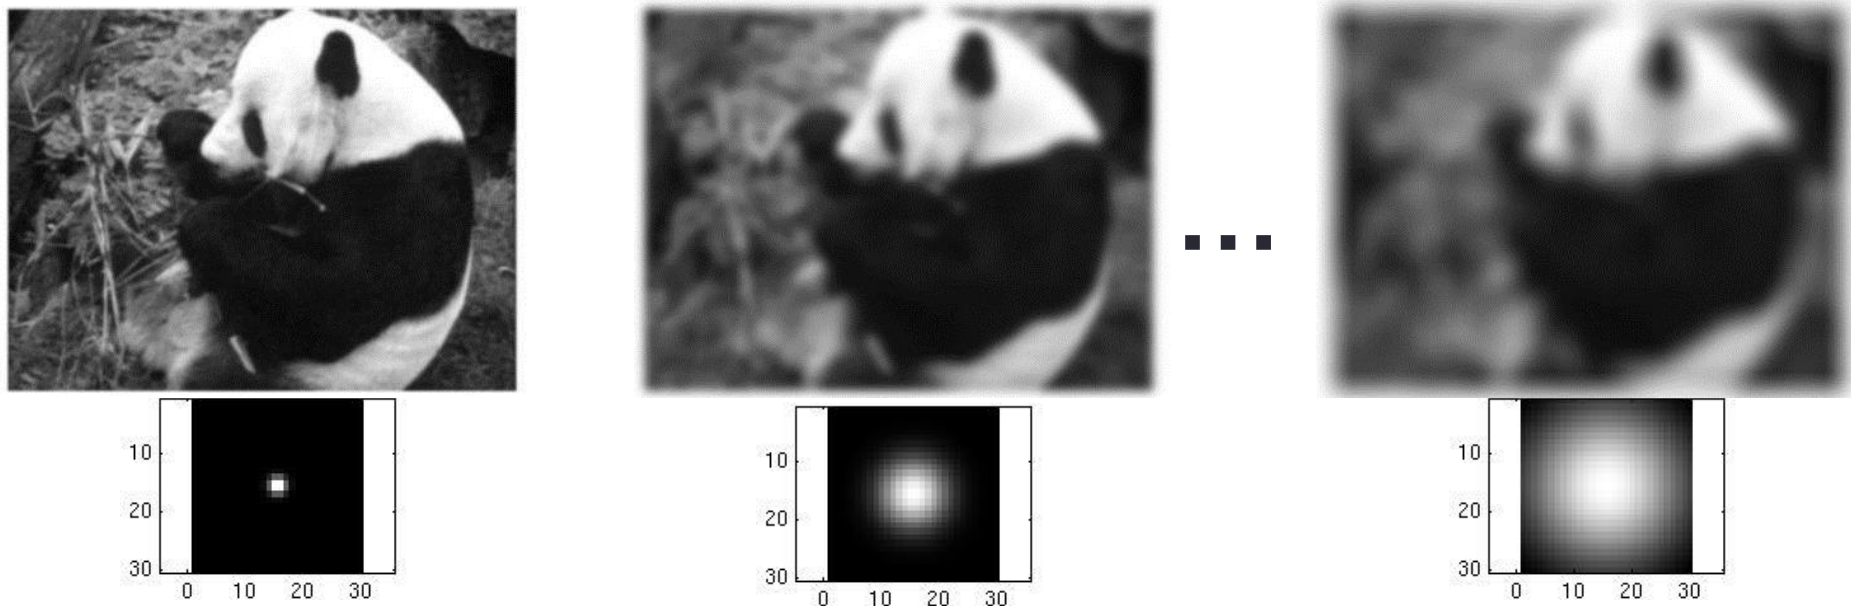
```
>> outim = imfilter(im, h) ;  
>> imshow(outim) ;
```



**outim**  
K. Grauman

# Smoothing with a Gaussian

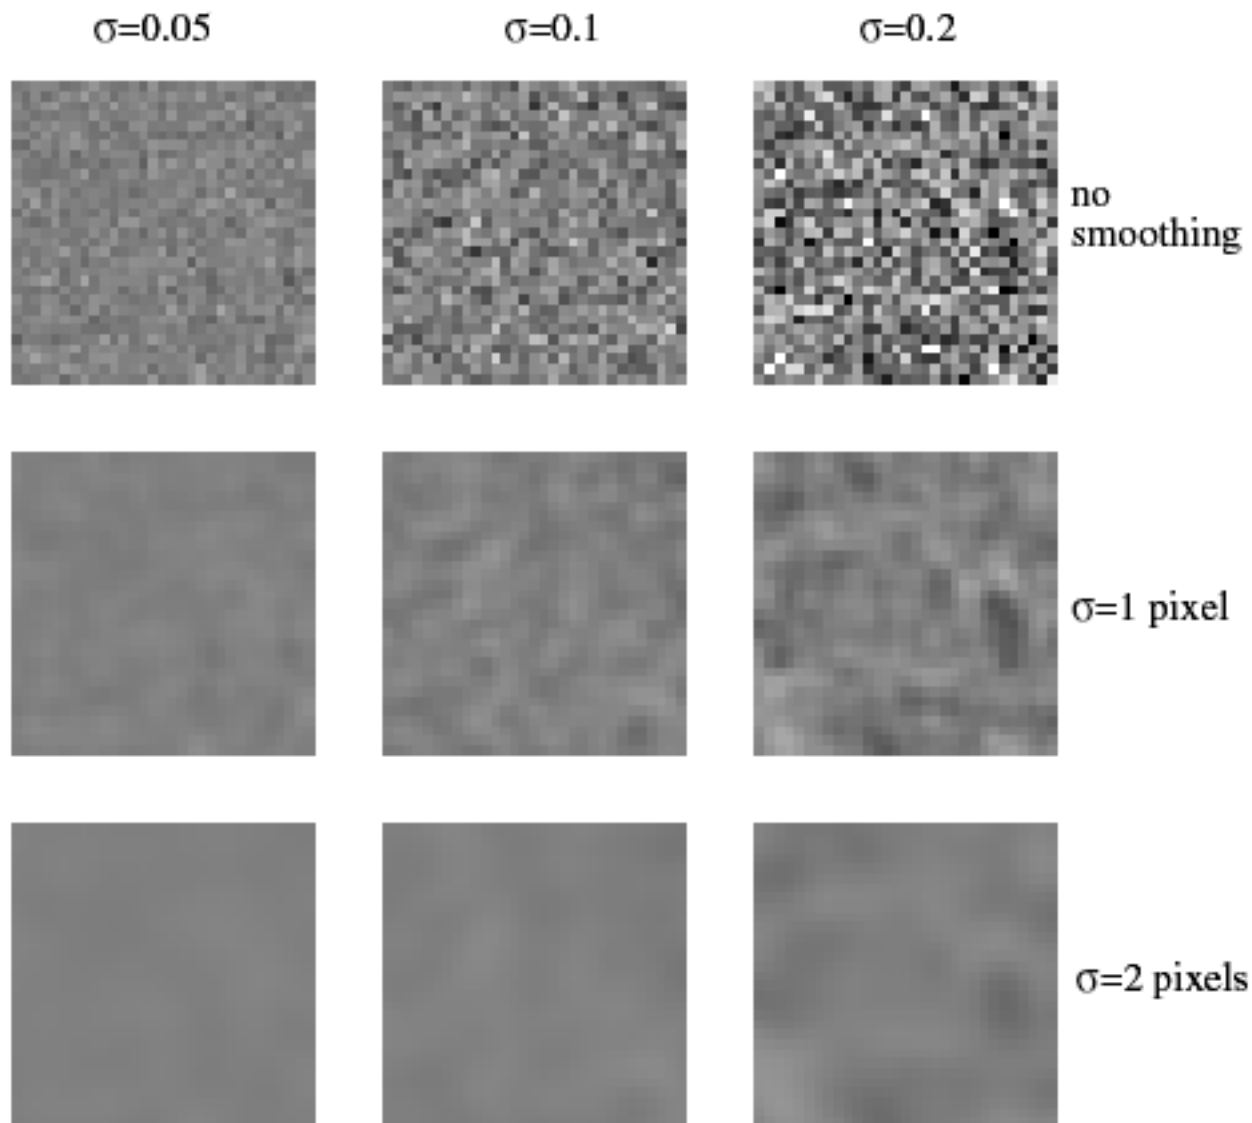
Parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



```
for sigma=1:3:10
    h = fspecial('gaussian', fsize, sigma);
    out = imfilter(im, h);
    imshow(out);
    pause;
end
```

# Keeping the two Gaussians straight...

More Gaussian noise (like earlier)  $\sigma \rightarrow$



Wider Gaussian smoothing kernel  $\sigma \rightarrow$

# And now some linear intuition...

An operator  $H$  (or system) is *linear* if two properties hold ( $f1$  and  $f2$  are some functions,  $a$  is a constant):

- **Superposition** (things sum):

$$H(f1 + f2) = H(f1) + H(f2) \quad (\text{looks like distributive law})$$


- **Scaling** (constant scales):

$$H(a \cdot f1) = a \cdot H(f1)$$


*Because it is sums and multiplies, the  
“filtering” operation we were doing are linear.*

# An impulse function...

- In the discrete world, an *impulse* is a very easy signal to understand: it's just a value of 1 at a single location.




- In the continuous world, an impulse is an idealized function that is very narrow and very tall so that it has a unit area. In the limit:



# An impulse response

- If I have an unknown system and I “put in” an impulse, the response is called the impulse response. (Duh?)




- So if the black box is linear you can describe  $H$  by  $h(x)$ . Why?




# Filtering an impulse signal


What is the result of filtering the impulse signal (image)  $F$  with the arbitrary kernel  $H$ ?




# Filtering an impulse signal




# Filtering an impulse signal




# Filtering an impulse signal




# Filtering an impulse signal




# Filtering an impulse signal



# Filtering an impulse signal




# Filtering an impulse signal







# Filtering an impulse signal



# Filtering an impulse signal



# “Filtering” an impulse signal



If you just “filter” meaning slide the kernel over the image you get a **reversed** response.

# Convolution

- Convolution:

- Flip where the filter is applied in both dimensions (bottom to top, right to left)
- Then apply cross-correlation

Centered at zero!




$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H * F$$



*Notation for  
convolution  
operator*



# One more thing...

- **Shift invariant:**
  - Operator behaves the same everywhere, i.e. the value of the output depends on the pattern in the image neighborhood, not the position of the neighborhood.

# Properties of convolution

- Linear & shift invariant

- Commutative:

$$f * g = g * f$$


- Associative

$$(f * g) * h = f * (g * h)$$

- Identity:

$$\text{unit impulse } e = [\dots, 0, 0, 1, 0, 0, \dots]. \quad f * e = f$$

- Differentiation:  $\frac{\partial}{\partial x} (f * g) = \frac{\partial f}{\partial x} * g$



We'll use  
this later!

# Convolution vs. correlation

## Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$

## Cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$


$$G = H \otimes F$$

For a Gaussian or box filter, how will the outputs differ?

If the input is an impulse signal, how will the outputs differ?

# Computational Complexity

- If an image is  $N \times N$  and a kernel (filter) is  $W \times W$ , how many multiplies do you need to compute a convolution?



- You need  $N \times N \times W \times W = N^2 W^2$ 
  - which can get big (ish)



# Separability

- In some cases, filter is separable, meaning you can get the square kernel by convolving a single column vector by some row vector:


$$\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

- To apply to an image you:
  - Convolve all rows
  - Convolve all resulting columns
- This used to be *very* important – instead of  $N \times N \times W \times W$  it's  $N \times N \times W \times 2$ 
  - So if your Kernel is a 31x31 filter you save a factor of 15


# Boundary issues

- What is the size of the output?
- Old MATLAB: `filter2(g, f, shape)`
  - *shape* = 'full': output size is sum of sizes of *f* and *g*
  - *shape* = 'same': output size is same as *f*
  - *shape* = 'valid': output size is difference of sizes of *f* and *g*


**full**



**same**



**valid**



# Boundary issues

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - **clip filter (black)**




# Boundary issues

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - **wrap around**



# Boundary issues

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - wrap around
    - **copy edge**



# Boundary issues


- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - wrap around
    - copy edge
    - **reflect across edge**





# Boundary issues

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods (new MATLAB):
    - clip filter (black): `imfilter(f, g, 0)`
    - wrap around: `imfilter(f, g, 'circular')`
    - copy edge: `imfilter(f, g, 'replicate')`
    - reflect across edge: `imfilter(f, g, 'symmetric')`

# Predict the filtered outputs


$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} = ?$$


$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix} = ?$$


$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} = ?$$



# Practice with linear filters



**Original**

0	0	0
0	1	0
0	0	0


?

# Practice with linear filters



**Original**

0	0	0
0	1	0
0	0	0



**Filtered  
(no change)**

# Practice with linear filters




**Original**

0	0	0
0	0	1
0	0	0

?

# Practice with linear filters



**Original**

0	0	0
0	0	1
0	0	0



*Center coordinate is 0,0!*



**Shifted left  
by 1 pixel  
with  
correlation**

# Practice with linear filters



**Original**

 $\frac{1}{9}$ 

1	1	1
1	1	1
1	1	1

?

# Practice with linear filters



**Original**

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1



**Blur (with a  
box filter)**

# Practice with linear filters



**Original**

0	0	0
0	2	0
0	0	0


–

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

# Practice with linear filters



**Original**

0	0	0
0	2	0
0	0	0

–

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1




**Sharpening filter**

**- Accentuates differences  
with local average**






# Filtering examples: sharpening



**before**



**after**

# Effect of smoothing filters

5x5

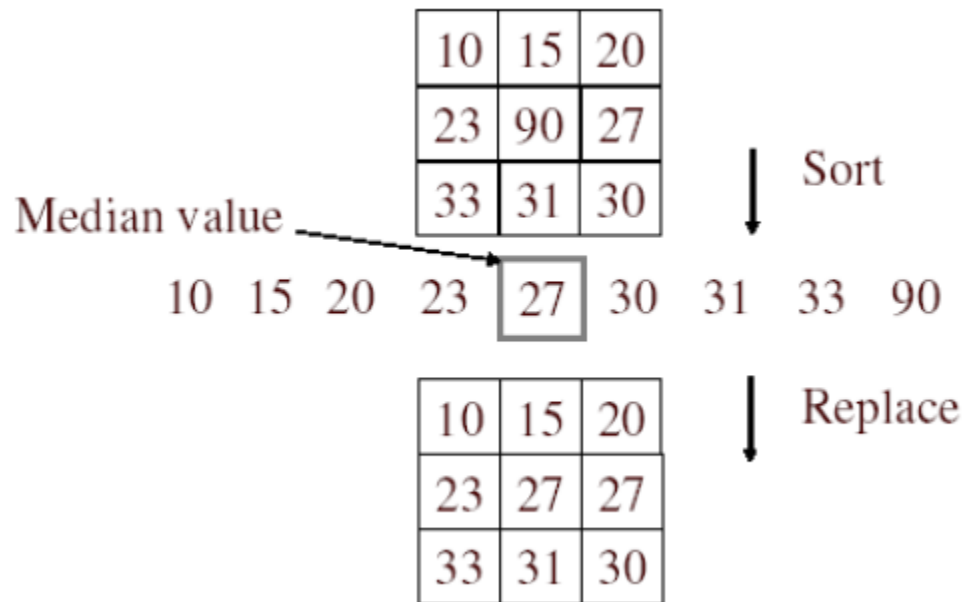


**Additive Gaussian noise**



**Salt and pepper noise**



# Median filter




- No new pixel values introduced
- Removes spikes: good for impulse, salt & pepper noise
- Linear?

# Median filter

**Salt and  
pepper  
noise** →




← **Median  
filtered**



**Plots of a row of the  
image**

# Median filter

- Median filter is edge preserving



# Exercise

- Design a kernel for:
  - Horizontal derivative
  - Vertical derivative
  - Second derivative
- What If:
  - H is the max/min operator?
  - We use an affine operator?

# Homework For Next Class

- Read slide set
- Prepare 3 questions
  - Option 1: on things that you don't understand
  - Option 2: on things that you think would make good exam questions
- Do it individually
- Write them down and cut them into 3 strips: 1 per questions and distribute them.
- **BRING ENOUGH LAPTOPS** to access the class.